CIS501 – Computer Architecture	Extra Practice
Prof. Martin	2011

Name: _____

This document contains some extra questions to give you some extra practice on some of the topics we discussed.

1. **Pipelines, Superscalar, and Scheduling I**. Below is the code for calculates the sum of the square of the difference of two arrays. The assembly is a simple RISC-like ISA that uses condition codes for branch outcomes. The BRn instruction means "branch if negative".

```
loop:
   LD    R4,   R1,   0    ; R4 <- Mem[R1+0]
   LD    R5,   R2,   0    ; R5 <- Mem[R2+0]
   SUB   R6,   R4,   R5   ; R6 <- R4 - R5
   MUL   R7,   R6,   R6   ; R7 <- R6 * R6
   ADD   R3,   R3,   R7   ; R3 <- R3 + R7
   ADD   R1,   R1,   1    ; R1 <- R1 + 1
   ADD   R2,   R2,   1    ; R2 <- R2 + 1
   ADD   R0,   R0,   1    ; R0 <- R0 + 1
   CMP   R0,   20
   BRn   loop
```

For the questions below, consider a 5-stage pipeline with a single cycle load-to-use penalty. All other instructions execute in a single cycle. All branches are predicted perfectly. For the superscalar pipelines assume any mix of instructions is allows in a given cycle (for example, no limit on the number of loads per cycle), and because the BRn reads the condition codes set by CMP, BRn and CMP are not allowed to execute in the same cycle.

(a) When executing the above code on a scalar (single-issue) pipeline, how many cycles per loop iteration? What is the CPI?

(b) When executing the above code on a two-way (dual-issue) superscalar pipeline, how many cycles per loop iteration? What is the CPI?

(c) Reschedule the code to achieve the best CPI on this dual-issue pipeline. Instead of rewriting the code, just fill in the cycle the instruction begins execution in the "____" in the code below.

```
loop:

    ____    LD    R4,   R1,   0     ; R4 <- Mem[R1+0]

    ____    LD    R5,   R2,   0     ; R5 <- Mem[R2+0]

    ____    SUB   R6,   R4,   R5    ; R6 <- R4 - R5

    ____    MUL   R7,   R6,   R6    ; R7 <- R6 * R6

    ____    ADD   R3,   R3,   R7    ; R3 <- R3 + R7

    ____    ADD   R1,   R1,   1     ; R1 <- R1 + 1

    ____    ADD   R2,   R2,   1     ; R2 <- R2 + 1

    ____    ADD   R0,   R0,   1     ; R0 <- R0 + 1

    ____    CMP   R0,   20

    ____    BRn   loop
```

(d) For this rescheduled code, how many cycles per loop iteration? What is the CPI?

2. **Pipelines, Superscalar, and Scheduling II**.

   Below is the same code as the previous question, but it has been loop unrolled once (two iterations converting into one):

```
loop:
  LD   R4,   R1,   0     ; R4 <- Mem[R1+0]
  LD   R5,   R2,   0     ; R5 <- Mem[R2+0]
  SUB  R6,   R4,   R5    ; R6 <- R4 - R5
  MUL  R7,   R6,   R6    ; R7 <- R6 * R6
  ADD  R3,   R3,   R7    ; R3 <- R3 + R7
  LD   R8,   R1,   1     ; R8 <- Mem[R1+1]
  LD   R9,   R2,   1     ; R9 <- Mem[R2+1]
  SUB  R10, R8,   R9    ; R10 <- R8 - R9
  MUL  R11, R10, R10   ; R11 <- R10 * R10
  ADD  R3,   R3,   R11   ; R3 <- R3 + R11
  ADD  R1,   R1,   2     ; R1 <- R1 + 2
  ADD  R2,   R2,   2     ; R2 <- R2 + 2
  ADD  R0,   R0,   2     ; R0 <- R0 + 2
  CMP  R0,   20
  BRn  loop
```

   (a) When executing the above code on a two-way (dual-issue) superscalar pipeline, how many cycles per loop iteration? What is the CPI?

   (b) Assume for a moment you have an infinite-issue width processor (it can execute an unbounded number if independent instructions per cycle). What is the best cycles per loop and CPI possible?

   (c) What is the minimum superscalar issue width necessary to achieve this best-case CPI?

(d) Show the rescheduled code that achieves the above best-base CPI. Instead of rewriting the code, just fill in the cycle the instruction begins execution in the "____" in the code below.

```
loop:

    ____    LD    R4,   R1,   0     ; R4 <- Mem[R1+0]

    ____    LD    R5,   R2,   0     ; R5 <- Mem[R2+0]

    ____    SUB   R6,   R4,   R5    ; R6 <- R4 - R5

    ____    MUL   R7,   R6,   R6    ; R7 <- R6 * R6

    ____    ADD   R3,   R3,   R7    ; R3 <- R3 + R7

    ____    LD    R8,   R1,   1     ; R8 <- Mem[R1+1]

    ____    LD    R9,   R2,   1     ; R9 <- Mem[R2+1]

    ____    SUB   R10,  R8,   R9    ; R10 <- R8 - R9

    ____    MUL   R11,  R10,  R10   ; R11 <- R10 * R10

    ____    ADD   R3,   R3,   R11   ; R3 <- R3 + R11

    ____    ADD   R1,   R1,   2     ; R1 <- R1 + 2

    ____    ADD   R2,   R2,   2     ; R2 <- R2 + 2

    ____    ADD   R0,   R0,   2     ; R0 <- R0 + 2

    ____    CMP   R0,   20

    ____    BRn   loop
```

To help you understand the operation of caches, the following questions give the current state of the caches and memory and a single byte-sized memory operation (load or store). That is, all addresses are are byte-granularity addresses.

Your task is to show us how the cache(s) and memory change when the instruction provided (load or store) is executed by crossing out the part that changed, and filling in what changed (for example, bring in the data block on a cache miss, set the tag, adjust the dirty bit, switch the LRU bit, write a dirty block back to memory, etc.).

For loads, also **circle the byte loaded by the instruction**. For stores, we supply both an address and a byte value to be written. Place that value into the cache, and **circle that byte as well**.

Each of the questions is independent (that is, what happened in the previous question has no impact on the next question).

All the caches are write-back, write-allocate caches. The block size in all cases is four bytes. As the size of the cache in the examples is larger than what we can fit on the paper, only the entries for the blocks in question are shown.

All of the addresses are given in *binary* (one bit per digit). To save space, all of the data values are given as hex values (4 bits per digit).

- For the single-core caches, a block can be in one of the following states:
    - C (*Clean*)
    - D (*Dirty*)

- For the set associative caches, if the *LRU* (least-recently used) field is either 0 (indicating that *Way 0* is has the LRU block) or 1 (indicating that *Way 1* has the LRU block).

- For the multi-core cache examples, it uses an MSI protocol where each block in the cache can be:
    - I (*Invalid*)
    - S (*Shared*, which is read-only and clean)
    - M (*Modified*, which is read-write and dirty)

    The memory also contains a state for each block:
    - I (*Idle*, which indicates no cores are caching the block)
    - S (*Shared*, which indicates at least one core is sharing the block and the memory is up to date)
    - M (*Modified*, which indicates one and only one core is caching the block in the Modified state; the memory is not necessarily up to date)

    In addition to the state, the memory has a *Sharers* field for tracking with processors (if any) are caching a particular block.

3. Direct Mapped Caches I [3 points]

```
load 1101010111010110
```

— Cache —

Index	Tag	Data 00	Data 01	Data 10	Data 11	State
⋮	⋮	⋮	⋮	⋮	⋮	⋮
0101110101	1101	BD	FE	AB	DE	D
⋮	⋮	⋮	⋮	⋮	⋮	⋮
1101011101	0101	CD	CA	CC	CB	C
⋮	⋮	⋮	⋮	⋮	⋮	⋮

— Memory —

Address	Data 00	Data 01	Data 10	Data 11
⋮	⋮	⋮	⋮	⋮
010101011110101xx	BB	BE	AE	FE
⋮	⋮	⋮	⋮	⋮
010111010111101xx	CD	CA	CC	CB
⋮	⋮	⋮	⋮	⋮
110101011110101xx	DE	DC	DD	DF
⋮	⋮	⋮	⋮	⋮
110111010111101xx	AD	BA	BE	CD
⋮	⋮	⋮	⋮	⋮

4. **Direct Mapped Caches II [3 points]**

```
store 0101110101110101 <- FF
```

— Cache —

Index	Tag	Data 00	Data 01	Data 10	Data 11	State
⋮	⋮	⋮	⋮	⋮	⋮	⋮
0101110101	1101	BD	FE	AB	DE	D
⋮	⋮	⋮	⋮	⋮	⋮	⋮
1101011101	0101	CD	CA	CC	CB	C
⋮	⋮	⋮	⋮	⋮	⋮	⋮

— Memory —

Address	Data 00	Data 01	Data 10	Data 11
⋮	⋮	⋮	⋮	⋮
010101011110101xx	BB	BE	AE	FE
⋮	⋮	⋮	⋮	⋮
010111010111101xx	CD	CA	CC	CB
⋮	⋮	⋮	⋮	⋮
110101011110101xx	DE	DC	DD	DF
⋮	⋮	⋮	⋮	⋮
110111010111101xx	AD	BA	BE	CD
⋮	⋮	⋮	⋮	⋮

5. Direct Mapped Caches III [3 points]

```
load 0101010111010110
```

— Cache —

Index	Tag	Data 00	Data 01	Data 10	Data 11	State
⋮	⋮	⋮	⋮	⋮	⋮	⋮
0101110101	1101	BD	FE	AB	DE	D
⋮	⋮	⋮	⋮	⋮	⋮	⋮
1101011101	0101	CD	CA	CC	CB	C
⋮	⋮	⋮	⋮	⋮	⋮	⋮

— Memory —

Address	Data 00	Data 01	Data 10	Data 11
⋮	⋮	⋮	⋮	⋮
010101011110101xx	BB	BE	AE	FE
⋮	⋮	⋮	⋮	⋮
010111010111101xx	CD	CA	CC	CB
⋮	⋮	⋮	⋮	⋮
110101011110101xx	DE	DC	DD	DF
⋮	⋮	⋮	⋮	⋮
110111010111101xx	AD	BA	BE	CD
⋮	⋮	⋮	⋮	⋮

6. **Direct Mapped Caches IV [3 points]**

```
store 1101110101110101 <- FF
```

— Cache —

Index	Tag	Data				State
		00	01	10	11	
⋮	⋮	⋮	⋮	⋮	⋮	⋮
0101110101	1101	BD	FE	AB	DE	D
⋮	⋮	⋮	⋮	⋮	⋮	⋮
1101011101	0101	CD	CA	CC	CB	C
⋮	⋮	⋮	⋮	⋮	⋮	⋮

— Memory —

Address	Data			
	00	01	10	11
⋮	⋮	⋮	⋮	⋮
01010101110101xx	BB	BE	AE	FE
⋮	⋮	⋮	⋮	⋮
01011101011101xx	CD	CA	CC	CB
⋮	⋮	⋮	⋮	⋮
11010101110101xx	DE	DC	DD	DF
⋮	⋮	⋮	⋮	⋮
11011101011101xx	AD	BA	BE	CD
⋮	⋮	⋮	⋮	⋮

7. **Set Associative Caches I [3 points]**

```
load 1101110101010000
```

— Cache —

Index	Way 0						LRU	Way 1					
	Tag	**Data**				**State**		**Tag**	**Data**				**State**
		00	**01**	**10**	**11**				**00**	**01**	**10**	**11**	
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
101010100	01011	AB	BD	EF	FA	C	1	11011	BD	BC	FE	DF	D
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮

— Memory —

Address	Data			
	00	**01**	**10**	**11**
⋮	⋮	⋮	⋮	⋮
010111010100xx	AB	BD	EF	FA
⋮	⋮	⋮	⋮	⋮
101011010100xx	CC	DD	FB	AD
⋮	⋮	⋮	⋮	⋮
110111010100xx	BF	CA	FE	FC
⋮	⋮	⋮	⋮	⋮

8. **Set Associative Caches II [3 points]**

```
load 1010110101010010
```

— Cache —

Index	Way 0 Tag	Way 0 Data 00	01	10	11	Way 0 State	LRU	Way 1 Tag	Way 1 Data 00	01	10	11	Way 1 State
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
101010100	01011	AB	BD	EF	FA	C	1	11011	BD	BC	FE	DF	D
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮

— Memory —

Address	Data 00	01	10	11
⋮	⋮	⋮	⋮	⋮
01011101010100xx	AB	BD	EF	FA
⋮	⋮	⋮	⋮	⋮
10101101010100xx	CC	DD	FB	AD
⋮	⋮	⋮	⋮	⋮
11011101010100xx	BF	CA	FE	FC
⋮	⋮	⋮	⋮	⋮

9. Multicore Caches I [3 points]

Processor 0: `load 1110010101011010`

— Processor 0's Cache —

Index	Tag	Data 00	Data 01	Data 10	Data 11	State
⋮	⋮	⋮	⋮	⋮	⋮	⋮
0101010110	0101	AF	AB	EC	EA	M
⋮	⋮	⋮	⋮	⋮	⋮	⋮
1101010110	0101	FA	BA	CE	AE	S
⋮	⋮	⋮	⋮	⋮	⋮	⋮

— Processor 1's Cache —

Tag	Data 00	Data 01	Data 10	Data 11	State
⋮	⋮	⋮	⋮	⋮	⋮
1101	DE	AD	BE	EF	S
⋮	⋮	⋮	⋮	⋮	⋮
0101	FA	BA	CE	AE	S
⋮	⋮	⋮	⋮	⋮	⋮

— Memory —

Address	Data 00	Data 01	Data 10	Data 11	State	Sharers
⋮	⋮	⋮	⋮	⋮	⋮	⋮
010101010110xx	DD	AA	DD	EE	M	P0
⋮	⋮	⋮	⋮	⋮	⋮	⋮
010111010110xx	FA	BA	CE	AE	S	P0, P1
⋮	⋮	⋮	⋮	⋮	⋮	⋮
110101010110xx	DE	AD	BE	EF	S	P1
⋮	⋮	⋮	⋮	⋮	⋮	⋮
111001010110xx	CE	ED	DE	EC	I	—
⋮	⋮	⋮	⋮	⋮	⋮	⋮

10. **Multicore Caches II [3 points]**

Processor **1**: `load 0101010101011000`

— Processor 0's Cache —

Index	Tag	Data				State
		00	01	10	11	
⋮	⋮	⋮	⋮	⋮	⋮	⋮
0101010110	0101	AF	AB	EC	EA	M
⋮	⋮	⋮	⋮	⋮	⋮	⋮
1101010110	0101	FA	BA	CE	AE	S
⋮	⋮	⋮	⋮	⋮	⋮	⋮

— Processor 1's Cache —

Tag	Data				State
	00	01	10	11	
⋮	⋮	⋮	⋮	⋮	⋮
1101	DE	AD	BE	EF	S
⋮	⋮	⋮	⋮	⋮	⋮
0101	FA	BA	CE	AE	S
⋮	⋮	⋮	⋮	⋮	⋮

— Memory —

Address	Data				State	Sharers
	00	01	10	11		
⋮	⋮	⋮	⋮	⋮	⋮	⋮
01010101010110xx	DD	AA	DD	EE	M	P0
⋮	⋮	⋮	⋮	⋮	⋮	⋮
01011101010110xx	FA	BA	CE	AE	S	P0, P1
⋮	⋮	⋮	⋮	⋮	⋮	⋮
11010101010110xx	DE	AD	BE	EF	S	P1
⋮	⋮	⋮	⋮	⋮	⋮	⋮
11100101010110xx	CE	ED	DE	EC	I	—
⋮	⋮	⋮	⋮	⋮	⋮	⋮

11. Multicore Caches III [3 points]

Processor **1**: `store 0101010101011001 <- FF`

— Processor 0's Cache —

Index	Tag	Data				State
		00	01	10	11	
⋮	⋮	⋮	⋮	⋮	⋮	⋮
0101010110	0101	AF	AB	EC	EA	M
⋮	⋮	⋮	⋮	⋮	⋮	⋮
1101010110	0101	FA	BA	CE	AE	S
⋮	⋮	⋮	⋮	⋮	⋮	⋮

— Processor 1's Cache —

Tag	Data				State
	00	01	10	11	
⋮	⋮	⋮	⋮	⋮	⋮
1101	DE	AD	BE	EF	S
⋮	⋮	⋮	⋮	⋮	⋮
0101	FA	BA	CE	AE	S
⋮	⋮	⋮	⋮	⋮	⋮

— Memory —

Address	Data				State	Sharers
	00	01	10	11		
⋮	⋮	⋮	⋮	⋮	⋮	⋮
01010101010110xx	DD	AA	DD	EE	M	P0
⋮	⋮	⋮	⋮	⋮	⋮	⋮
01011101010110xx	FA	BA	CE	AE	S	P0, P1
⋮	⋮	⋮	⋮	⋮	⋮	⋮
11010101010110xx	DE	AD	BE	EF	S	P1
⋮	⋮	⋮	⋮	⋮	⋮	⋮
11100101010110xx	CE	ED	DE	EC	I	—
⋮	⋮	⋮	⋮	⋮	⋮	⋮

12. **Multicore Caches IV [3 points]**

Processor 0: `store 1101010101011010 <- FF`

— Processor 0's Cache —

Index	Tag	Data 00	Data 01	Data 10	Data 11	State
⋮	⋮	⋮	⋮	⋮	⋮	⋮
0101010110	0101	AF	AB	EC	EA	M
⋮	⋮	⋮	⋮	⋮	⋮	⋮
1101010110	0101	FA	BA	CE	AE	S
⋮	⋮	⋮	⋮	⋮	⋮	⋮

— Processor 1's Cache —

Tag	Data 00	Data 01	Data 10	Data 11	State
⋮	⋮	⋮	⋮	⋮	⋮
1101	DE	AD	BE	EF	S
⋮	⋮	⋮	⋮	⋮	⋮
0101	FA	BA	CE	AE	S
⋮	⋮	⋮	⋮	⋮	⋮

— Memory —

Address	Data 00	Data 01	Data 10	Data 11	State	Sharers
⋮	⋮	⋮	⋮	⋮	⋮	⋮
010101010101 10xx	DD	AA	DD	EE	M	P0
⋮	⋮	⋮	⋮	⋮	⋮	⋮
010111010101 10xx	FA	BA	CE	AE	S	P0, P1
⋮	⋮	⋮	⋮	⋮	⋮	⋮
110101010101 10xx	DE	AD	BE	EF	S	P1
⋮	⋮	⋮	⋮	⋮	⋮	⋮
111001010101 10xx	CE	ED	DE	EC	I	—
⋮	⋮	⋮	⋮	⋮	⋮	⋮

13. Multicore Caches V [3 points]

Processor 0: `store 0101110101011011 <- FF`

— Processor 0's Cache —

Index	Tag	Data				State
		00	01	10	11	
⋮	⋮	⋮	⋮	⋮	⋮	⋮
0101010110	0101	AF	AB	EC	EA	M
⋮	⋮	⋮	⋮	⋮	⋮	⋮
1101010110	0101	FA	BA	CE	AE	S
⋮	⋮	⋮	⋮	⋮	⋮	⋮

— Processor 1's Cache —

Tag	Data				State
	00	01	10	11	
⋮	⋮	⋮	⋮	⋮	⋮
1101	DE	AD	BE	EF	S
⋮	⋮	⋮	⋮	⋮	⋮
0101	FA	BA	CE	AE	S
⋮	⋮	⋮	⋮	⋮	⋮

— Memory —

Address	Data				State	Sharers
	00	01	10	11		
⋮	⋮	⋮	⋮	⋮	⋮	⋮
01010101010110xx	DD	AA	DD	EE	M	P0
⋮	⋮	⋮	⋮	⋮	⋮	⋮
01011101010110xx	FA	BA	CE	AE	S	P0, P1
⋮	⋮	⋮	⋮	⋮	⋮	⋮
11010101010110xx	DE	AD	BE	EF	S	P1
⋮	⋮	⋮	⋮	⋮	⋮	⋮
11100101010110xx	CE	ED	DE	EC	I	—
⋮	⋮	⋮	⋮	⋮	⋮	⋮