

Question 1

The algorithm is for each node v and its left child v' and right child v'' , if $l_{(v,v')} \neq l_{(v,v'')}$, then adjust the lengths of two edges such that let $l_{(v,v')} = l_{(v,v'')} = \max\{l_{(v,v')}, l_{(v,v'')}\}$. After we visit each node by using BFS and adjust corresponding edges, the resulting binary tree has zero skew and the total edge length is as small as possible.

Observation 1: For any subtree rooted at v with all its descendent leaves, x_1 to x_k , after we make every path from v to x_i , where $i \in [1..k]$, equal, there always is a path from v to a leaf x_m , none of whose edges will be increased, for the case that the total length is as small as possible,

Proof. Suppose not. Let l_i denote the length of the path from v to x_i , where $i \in [1..k]$. Suppose we add $\Delta_i > 0$ to each path from v to x_i , such that all $l_i + \Delta_i$ are equal and the total length is as small as possible. But we can find the smallest Δ_m and subtract Δ_m from each Δ_i , such that for each path from v to x_i , $(l_i + \Delta_i - \Delta_m)$ is still equal to each other and get the a smaller total length. Additionally, because $\Delta_m = 0$, the path from v to x_m is untouched, therefore no edges along the path from v to x_m is increased. \square

Proof. Suppose not. Then there are other solutions better than the algorithm provided. Let v be any node that do not follow the above algorithm. There are three alternate solutions that can be better,

- 1) we let $l_{(v,v')} = l_{(v,v'')} = l' > \max\{l_{(v,v')}, l_{(v,v'')}\}$. Then at node v , all paths from v are increased, which contradicts the observation 1.
- 2) we let $l_{(v,v')} > l_{(v,v'')}$ after we change the lengths of the corresponding edges. Therefore, the length of each path from v'' to a leaf has to be increased by $l_{(v,v')} - l_{(v,v'')}$, in order to keep zero skew from root to leaves. Then all paths from v'' are increased, which contradicts the observation 1.
- 3) we let $l_{(v,v')} < l_{(v,v'')}$ after we change the lengths of the corresponding edges. Symmetrically, we will get a contradiction.

In sum, any alternate solution is no better than the provided one. So the provided algorithm is optimal. \square

The BFS takes $O(n)$ for the tree traverse, and at each node the operation is constant time, therefore, the total runtime is $O(n)$.

Question 2

First, we make all weights of X-edges larger than Y-edges, and use Prim's algorithm to find a minimum spanning T_1 with least number of X-edges x_{min} . Second, we make all weights of X-edges smaller than Y-edges, and use Prim's algorithm to find a minimum spanning tree T_2 with largest number of X-edges x_{max} .

If $k < x_{min}$ or $k > x_{max}$, we can conclude that no such tree that with exactly k edges labeled X exists. If $k = x_{min} = x_{max}$, then the minimum spanning tree produced above is the spanning tree with exactly k edges labeled X . For the case that $x_{min} < k < x_{max}$, we need to find a tree with exactly k edges labeled X by swapping some edges of the minimum spanning tree, T_1 or T_2 produced above. The swapping method is described below.

We choose to start from the spanning tree T_1 with x_{min} edges. First, we pick a edge $e \in T_2 - T_1$ and the graph $T_1 \cup \{e\}$ contains a cycle C . Then the cycle must contain a edge $e' \notin T_2$. We form a new graph $T'_1 = T_1 \cup \{e\} - \{e'\}$ and T'_1 has at least one edge belongs to T_2 . Because T_1 can have at most $n - 1$ edges different from T_2 , namely T_1 and T_2 have no common edge, therefore, after at most $n - 1$ step mentioned above, T_1 will become T_2 , and the number of X-edges in the spanning tree will be increased from x_{min} to x_{max} . In the process, each time only one edge is swapped, therefore the number of X-edges in the spanning tree will increase at most 1. As a result, we can get a exactly k X-edges spanning tree during the swapping process.

The minimum spanning tree algorithm will take $O(m \log n)$. Each swapping will take $O(m)$, and the whole swapping will take $O(mn)$. Therefore the total runtime will be $O(mn)$

Question 3

a) First, the general Resource Reservation Problem is NP, because, given any allocation schema, by checking each process whether its request is satisfied, we can tell whether there are k active processes.

Second, we use independent set problem, and if we can show $IS \leq_P RRP$, then we can conclude that the general Resource Reservation Problem is NP-hard.

Given any arbitrary instance of independent set problem, where the graph is $G = (V, E)$ and the number of independent set is k , we will first construct an equivalent instance of resource reservation problem. We let each node $v \in V$ denote one process, and there is an edge (u, v) if process u and process v both request at least one common resource. If we can find a k -sized independent set in G , then the corresponding nodes are the active processes in resource reservation problem.

Proof. First, we prove " \Rightarrow ". When we have a k -sized independent set, we have k nodes are disjoint, namely k processes request no common resources. Therefore, these k processes can be active.

Second, we prove " \Leftarrow ". When we have k active processes, these processes will request no common resources at all. Then, in the corresponding graph G , these k nodes will have no edge

between any two of them. Therefore, these k nodes form a k -sized independent set. \square

Hereby, we have resource reservation problem is both NP and NP-hard, then we can conclude that resource reservation problem is NP-complete.

b)

c)

d)

Question 4